

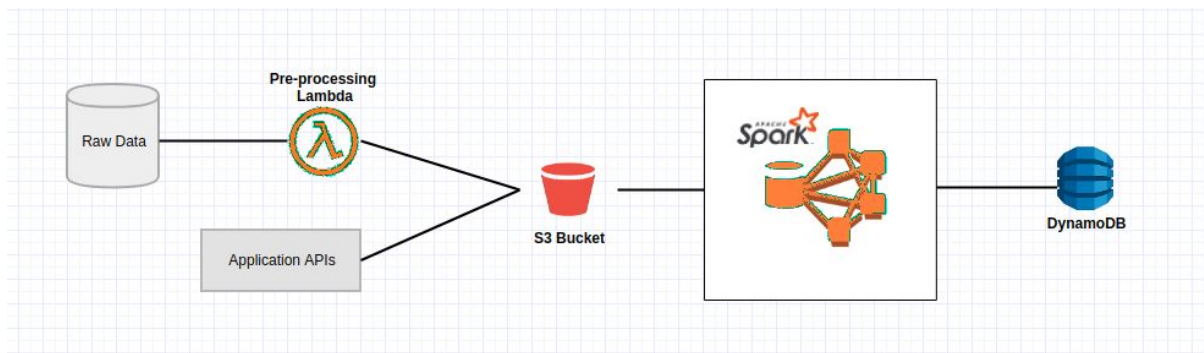
CHECKPOINT 2

DATA CENTER SCALE COMPUTING

PROBLEM DESCRIPTION

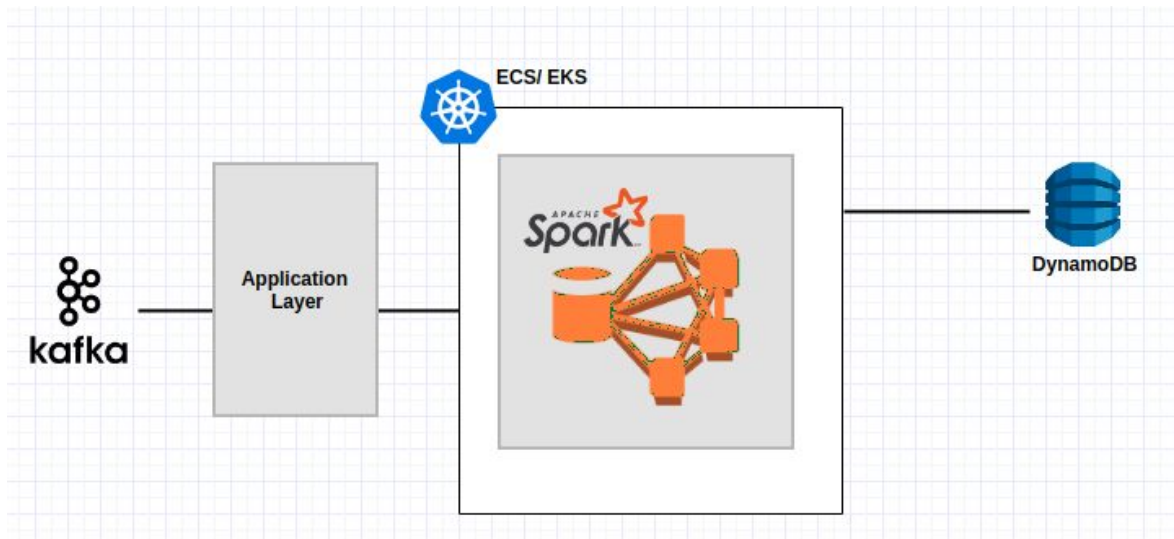
- Through this project, we aim to focus on scalability. In an application like Uber or any taxi service app, it is pertinent that the user requests of a cab are met in an acceptable amount of time. If the application delays a user's response a lot or fails to return a response, this could lead to a reduced popularity for the application.
- The application should be able to handle a large number of requests at the same time, to promote user satisfaction and also for monetary purposes.
- We plan to explore Kubernetes/Docker services (EKS/ECS) to handle multiple requests parallelly. We also plan to split independent requirements of a request and delegate them to different microservices so that they are computed parallelly and independently.

HIGH-LEVEL ARCHITECTURE



Pre-Processing data and storing the structured data in DynamoDB database

- Raw data would be put into some kind of database (Maybe S3)
- Pre-processing lambda function takes data from this database, preprocesses it and stores it in S3.
- Spark then takes care of creating tables and columns in DynamoDB.



Kafka to process “real time” requests

- Kafka would contain user requests. For example, a request to search for a cab at a particular location.
- This request would then go to the application layer which would structure the user query, maybe add additional context and send it to a node with Spark capabilities in EKS/ECS (Elastic Kubernetes Service or Elastic Container Service).
- EKS is Elastic Kubernetes Service which provides an option to configure a Kubernetes cluster. ECS is Elastic Container Service which provides an option of creating multiple dockers.
- We then configure some nodes in EKS/ECS to have Spark functionalities.
- When a request reaches the EKS/ECS a new node might be created on the fly or use an existing node with spark functionalities.
- Spark takes care of interacting with the DynamoDb and querying for results which it sends back to the application layer.
- Extended Goal: Show some kind of visualization of the result.

DATASET

- We will be using NYC taxi trip data from Yellow Taxi, Green Taxi and Uber from 2014 - [database](#)
- Description:
 - Format: .csv
 - The data requires minimal preprocessing. (Removal of unnecessary metadata)
 - The dataset is static.

- The data can be exported as a CSV. There is no need of a developer account.
- The data will be stored in S3 buckets.
- Sample dataset - (Next Page)

Column Name	Description	Type	
vendorid	A code indicating the TPEP provider that provided the reco...	Plain Text	T
pickup_datetime	The date and time when the meter was engaged.	Date & Time	
dropoff_datetime	The date and time when the meter was disengaged.	Date & Time	
Store_and_fwd_flag	This flag indicates whether the trip record was held in vehi...	Plain Text	T
rate_code	The final rate code in effect at the end of the trip. 1= Stand...	Number	#
Pickup_longitude		Number	#
Pickup_latitude	Latitude where the meter was engaged.	Number	#
Dropoff_longitude	Longitude where the meter was disengaged.	Number	#
Dropoff_latitude	Latitude where the meter was disengaged.	Number	#
Passenger_count	The number of passengers in the vehicle. This is a driver-e...	Number	#

Some Columns that are present in the [database](#)

CHALLENGES

- One of the challenges is to process the incoming requests fast and gain a major boost in the processing time through our architecture.
- Our architecture consists of multiple technologies like Kafka, EKS/ECS, ElasticSearch alongwith Microservices. The biggest challenge would be to integrate all the components and run the queries.
- There are technologies like Kubernetes, ECS, EKS, Kafka which our group is not familiar with. It would be a challenge to gain holistic conceptual understanding and recognize appropriate implementation strategies.
- Our metrics of measuring performance of our project would be to calculate the average running times of 100 different requests which requires a considerable processing power and time and compare the running times with and without using ECS/EKS etc.

GENERAL TASKS AND TIMELINE

Given that we are to present two checkpoints for the given project, we have divided our timeline accordingly.

Weeks	Tasks planned
Oct 23 - Oct 30	Project Proposal <ul style="list-style-type: none">● Teaming up to discuss architecture and write proposal
Nov 1 - Nov 13	<ul style="list-style-type: none">● Finalise project Design● Start setting up EKS/ECS● Preprocessing, data cleaning and structuring● Setting up Kafka● Designing the Application Layer
Nov 14- Nov 29	Project Implementation
	<ul style="list-style-type: none">● Integrating Spark with EKS/ECS● Structuring the queries from Spark on DynamoDB● Kafka processing● Kafka-Application Layer integration● Application Layer● Application Layer-EKS/ECS Integration● Writing Unit test cases
Nov 30 - Dec 10	<ul style="list-style-type: none">● Project Completion● Documentation● Integration testing● Performance Measure

TASK DIVISION AND TIMELINE

S. No.	Name	Tasks	Timeline
1	Akriti Kapur	<ul style="list-style-type: none">• Defining Architecture• Overlooking preprocessing, data cleaning and structuring• Kafka processing• Structure the queries• Creating test cases to produce “real-time” data	<ul style="list-style-type: none">• Week 0• Week 1• Week 2• Week 2• Week 3
2	Amith Gopal	<ul style="list-style-type: none">• Identifying Challenges• Containerization using EKS/ECS• Integration of Spark with EKS/ECS• Integration of Application Layer with EKS/ECS• Simple queries to test Spark-dynamoDB	<ul style="list-style-type: none">• Week 0• Week 1• Week 2• Week 3• Week 3
3	Sowmya	<ul style="list-style-type: none">• Identifying the AWS services• Containerization using EKS/ECS• Integration of Spark with EKS/ECS• Integration of Application Layer with EKS/ECS• Documentation	<ul style="list-style-type: none">• Week 0• Week 1• Week 2• Week 3• Week 3
4	Tarunianand	<ul style="list-style-type: none">• Finding datasets• Working with Akriti on Kafka processing• Developing application layer (Django etc.)• Creating test cases to produce “real-time” data• Documentation and Demonstration prep	<ul style="list-style-type: none">• Week 0• Week 1• Week 2• Week 3• Week 3

CONCERN

- Since our project is heavily dependent on the AWS services, we are concerned about the potential cost incurred by the time we complete the project.

CHECKPOINT 1

Changes to the existing proposal

1. Instead of using Dynamodb as the database, we would be using MongoDB database.

New Timeline

Work done so far:

Summary of the work done

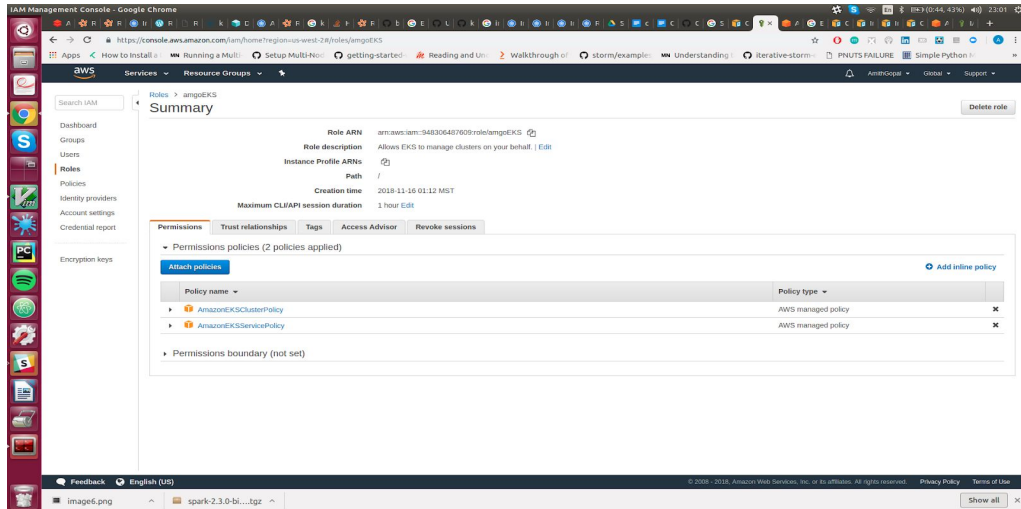
1. **Cleaning and Preprocessing of data**
2. **Storing the cleaned data in a database (MongoDB)**
3. **Setting up Kubernetes Cluster**
4. **Running a sample Spark job on the Kubernetes Cluster to verify the EKS setup**
5. **Setting up Minikube**
6. **Setting up a basic Django App**

Akriti Kapur:

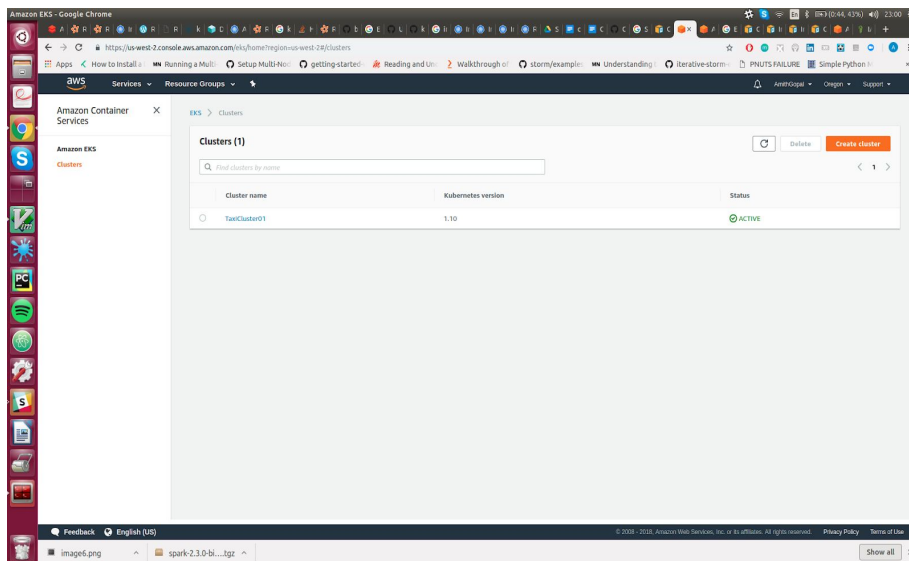
- Preprocessing architecture implemented
 - Lambda function to get rid of certain rows
 - Spark to operate on the sanitized data, create dataframe from the csv
 - Store the spark dataframe to a mongodb database
- Project Structured
 - Project settings, dependencies, environment variables added.

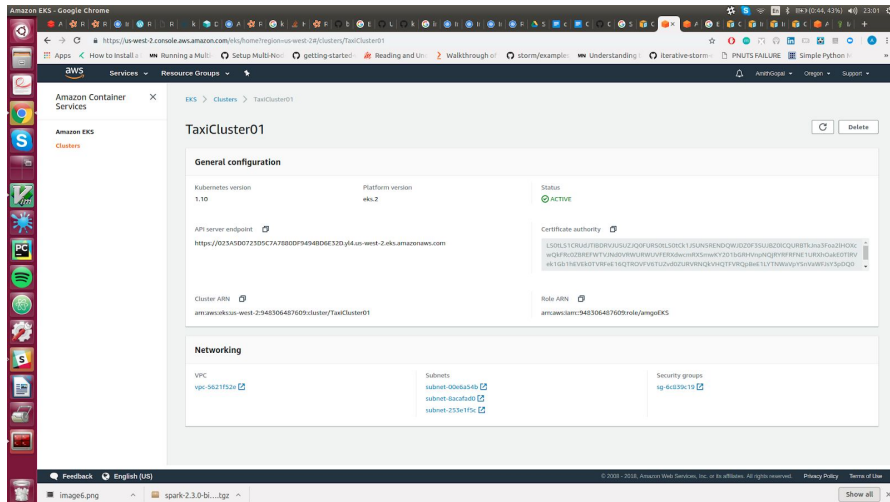
Amith Gopal

- I went through the tutorials, did research on prerequisites required before setting up EKS. After this, I installed kubectl, awscli and aws-iam-authenticator before proceeding to create EKS.
- I created the IAM roles, VPCs required for the creation of the EKS cluster.

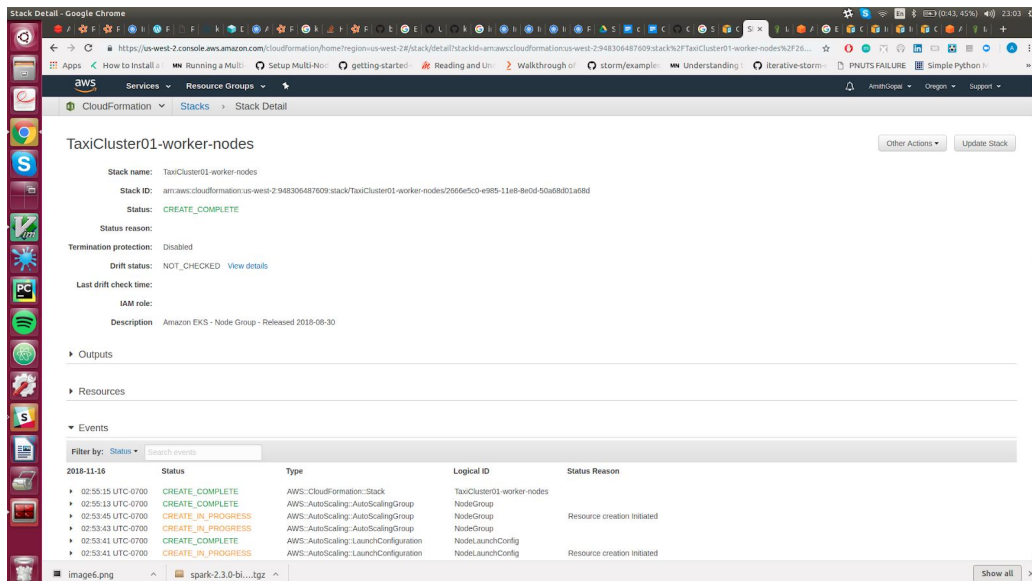


- I set up the EKS cluster which is the EKS control plane after the prerequisites.





- Next, I set up the worker nodes after the cluster was created and later configured kubectl with a .yml file to include the worker nodes.




```
/bin/bash
smith@smith-carbon: ~/Downloads/spark-2.3.0-bin-hadoop2.7$ kubectl get nodes --watch
NAME                                STATUS    ROLES    AGE   VERSION
ip-172-31-15-14.us-west-2.compute.internal Ready    <none>   20h   v1.10.3
ip-172-31-28-203.us-west-2.compute.internal Ready    <none>   20h   v1.10.3
ip-172-31-45-137.us-west-2.compute.internal Ready    <none>   20h   v1.10.3
ip-172-31-15-14.us-west-2.compute.internal Ready    <none>   20h   v1.10.3
ip-172-31-28-203.us-west-2.compute.internal Ready    <none>   20h   v1.10.3
ip-172-31-45-137.us-west-2.compute.internal Ready    <none>   20h   v1.10.3
```

- Lastly, I downloaded spark2.3 and ran a spark job on the kubernetes cluster using the command shown in the “Steps-followed-for-setting-EKS.txt” and verified the working of the kubernetes cluster.

```
/bin/bash
smith@smith-carbon: ~$ spark-submit --master k8s://https://223A5D0723D5C7A7880DF9494BD6E32D.y4.us-west-2.eks.amazonaws.com --deploy-mode cluster
--name spark-pi --class org.apache.spark.examples.SparkPi --conf spark.executor.instances=3 --conf spark.kubernetes.authenticate.driver.serviceAccountName=arn:aws:eks:us-west-2:948306487609:cluster/TaxiCluster01 --conf spark.kubernetes.container.image=kubespark/spark-init:v2.2.0-kubernetes-0.5.0 --conf spark.kubernetes.namespace=kube-system local:///home/smith/Downloads/spark-2.3.0-bin-hadoop2.7/examples/jars/spark-examples_2.11-2.3.0.jar
2018-11-16 23:29:49 WARN Utils:66 - Your hostname, smith-carbon resolves to a loopback address: 127.0.0.1; using 10.0.0.83 instead (on interface wipso)
2018-11-16 23:29:51 INFO LoggingPodStatusWatcherImpl:54 - State changed, new state:
pod name: spark-pi-0e809f6757a0381893f119f0b8fdcb6-driver
namespace: kube-system
labels: spark-app-selector -> spark-4fbb2d4d75f4c7d889feeb0d533abd4, spark-role -> driver
pod uid: 304be9a6-ea32-11e8-85d8-0a070ea86608
creation time: 2018-11-17T06:29:51Z
service account name: default
volumes: default-token-xf4pd
node name: N/A
start time: N/A
container images: N/A
phase: Pending
status: {}
2018-11-16 23:29:51 INFO LoggingPodStatusWatcherImpl:54 - State changed, new state:
pod name: spark-pi-0e809f6757a0381893f119f0b8fdcb6-driver
namespace: kube-system
labels: spark-app-selector -> spark-4fbb2d4d75f4c7d889feeb0d533abd4, spark-role -> driver
pod uid: 304be9a6-ea32-11e8-85d8-0a070ea86608
creation time: 2018-11-17T06:29:51Z
service account name: default
volumes: default-token-xf4pd
node name: ip-172-31-15-14.us-west-2.compute.internal
start time: N/A
container images: N/A
phase: Pending
status: {}
2018-11-16 23:29:51 INFO LoggingPodStatusWatcherImpl:54 - State changed, new state:
pod name: spark-pi-0e809f6757a0381893f119f0b8fdcb6-driver
namespace: kube-system
labels: spark-app-selector -> spark-4fbb2d4d75f4c7d889feeb0d533abd4, spark-role -> driver
pod uid: 304be9a6-ea32-11e8-85d8-0a070ea86608
creation time: 2018-11-17T06:29:51Z
service account name: default
volumes: default-token-xf4pd
node name: ip-172-31-15-14.us-west-2.compute.internal
start time: 2018-11-17T06:29:51Z
container images: kubespark/spark-init:v2.2.0-kubernetes-0.5.0
phase: Pending
status: {ContainerStatus(containerID=null, image=kubespark/spark-init:v2.2.0-kubernetes-0.5.0, imageID=, lastState=ContainerState(running=null, terminated=null, waiting=null, additionalProperties={}), name=spark-kubernetes-driver, ready=false, restartCount=0, state=ContainerState(running=null, terminated=null, waiting=ContainerStateWaiting(message=null, reason=ContainerCreating, additionalProperties={}), additionalProperties={}), additionalProperties={})}
2018-11-16 23:29:51 INFO Client:54 - Waiting for application spark-pi to finish...
2018-11-16 23:29:52 INFO LoggingPodStatusWatcherImpl:54 - State changed, new state:
pod name: spark-pi-0e809f6757a0381893f119f0b8fdcb6-driver
namespace: kube-system
labels: spark-app-selector -> spark-4fbb2d4d75f4c7d889feeb0d533abd4, spark-role -> driver
pod uid: 304be9a6-ea32-11e8-85d8-0a070ea86608
creation time: 2018-11-17T06:29:51Z
service account name: default
volumes: default-token-xf4pd
node name: ip-172-31-15-14.us-west-2.compute.internal
```

```
/bin/bash
pod name: spark-pi-0e809f6757a0381893f119f0b8fdcb6-driver
namespace: kube-system
labels: spark-app-selector -> spark-4fbb29ddd75f4c7d889feeb0d533abd4, spark-role -> driver
pod uid: 304be96f-ea32-11e8-85d8-0a070ea86608
creation time: 2018-11-17T06:29:51Z
service account name: default
volumes: default-token-xf4pd
node name: ip-172-31-15-14.us-west-2.compute.internal
start time: 2018-11-17T06:29:51Z
container images: kubernetes/spark-init:v2.2.0-kubernetes-0.5.0
phase: Pending
status: [ContainerStatus(containerID=null, image=kubernetes/spark-init:v2.2.0-kubernetes-0.5.0, imageID=, lastState=ContainerState(running=null,
terminated=null, waiting=null, additionalProperties={}), name=spark-kubernetes-driver, ready=false, restartCount=0, state=ContainerState(running=null, te
rminated=null, waiting=ContainerStateWaiting(message=null, reason=ContainerCreating, additionalProperties={}), additionalProperties={}), terminated=null, wait
ing=null)]
2018-11-16 23:29:51 INFO Client:54 - Waiting for application spark-pi to finish...
2018-11-16 23:29:52 INFO LoggingPodStatusWatcherImpl:54 - State changed, new state:
pod name: spark-pi-0e809f6757a0381893f119f0b8fdcb6-driver
namespace: kube-system
labels: spark-app-selector -> spark-4fbb29ddd75f4c7d889feeb0d533abd4, spark-role -> driver
pod uid: 304be96f-ea32-11e8-85d8-0a070ea86608
creation time: 2018-11-17T06:29:51Z
service account name: default
volumes: default-token-xf4pd
node name: ip-172-31-15-14.us-west-2.compute.internal
start time: 2018-11-17T06:29:51Z
container images: kubernetes/spark-init:v2.2.0-kubernetes-0.5.0
phase: Running
status: [ContainerStatus(containerID=docker://03b201be31bedc400538278faa8a17425540d889b4d4f04d3db9226c75b7404, image=kubernetes/spark-init:v2.2.
0-kubernetes-0.5.0, imageID=docker-pullable://kubernetes/spark-init@sha256:0fbec6a33d5c7293266ac6be7544e4ab595eea20d5eb8ab920f7c7dbd09ac2cd, lastState=Con
tainerState(running=null, terminated=null, waiting=null, additionalProperties={}), name=spark-kubernetes-driver, ready=true, restartCount=0, state=Contai
nerState(running=ContainerStateRunning(startedAt=Time(time=2018-11-17T06:29:52Z, additionalProperties={}), additionalProperties={}), terminated=null, wa
iting=null, additionalProperties={}), additionalProperties={})]
2018-11-16 23:29:53 INFO LoggingPodStatusWatcherImpl:54 - State changed, new state:
pod name: spark-pi-0e809f6757a0381893f119f0b8fdcb6-driver
namespace: kube-system
labels: spark-app-selector -> spark-4fbb29ddd75f4c7d889feeb0d533abd4, spark-role -> driver
pod uid: 304be96f-ea32-11e8-85d8-0a070ea86608
creation time: 2018-11-17T06:29:51Z
service account name: default
volumes: default-token-xf4pd
node name: ip-172-31-15-14.us-west-2.compute.internal
start time: 2018-11-17T06:29:51Z
container images: kubernetes/spark-init:v2.2.0-kubernetes-0.5.0
phase: Failed
status: [ContainerStatus(containerID=docker://03b201be31bedc400538278faa8a17425540d889b4d4f04d3db9226c75b7404, image=kubernetes/spark-init:v2.2.
0-kubernetes-0.5.0, imageID=docker-pullable://kubernetes/spark-init@sha256:0fbec6a33d5c7293266ac6be7544e4ab595eea20d5eb8ab920f7c7dbd09ac2cd, lastState=Con
tainerState(running=null, terminated=null, waiting=null, additionalProperties={}), name=spark-kubernetes-driver, ready=false, restartCount=0, state=Contai
nerState(running=null, terminated=ContainerStateTerminated(containerID=docker://03b201be31bedc400538278faa8a17425540d889b4d4f04d3db9226c75b7404, exitco
de=1, finishedAt=Time(time=2018-11-17T06:29:53Z, additionalProperties={}), message=null, reason=Error, signal=null, startedAt=Time(time=2018-11-17T06:29:
52Z, additionalProperties={}), additionalProperties={}), waiting=null, additionalProperties={}), additionalProperties={})]
2018-11-16 23:29:53 INFO LoggingPodStatusWatcherImpl:54 - Container final statuses:
Container name: spark-kubernetes-driver
Container image: kubernetes/spark-init:v2.2.0-kubernetes-0.5.0
Container state: Terminated
Exit code: 1
2018-11-16 23:29:53 INFO Client:54 - Application spark-pi finished
2018-11-16 23:29:53 INFO ShutdownHookManager:54 - Shutdown hook called
2018-11-16 23:29:53 INFO ShutdownHookManager:54 - Deleting directory /tmp/spark-9d1f3684-97a2-4693-a871-bdaa1c241f41
amith@ec2-carbon:~$ spark-submit --https://023a5d0723d5c7a7880df94948bd632d.y14.us-west-2.eks.amazonaws.com --deploy-mode cluster --name spark-pi --class
com.spark.kubernetes.namespace=kube-system local:///home/amith/Downloads/spark-2.3.0-bin-hadoop2.7/examples/jars/spark-examples_2.11-2.3.0.jar
```

- I learnt more about Spark's integration with EKS and other functionalities and properties of EKS which can be used

Sowmya Ramakrishnan :

- Learnt Kubernetes basics, concepts, commands, creating pods and clusters with labels and selectors, scaling up/ down, zero downtime deployment.(Documentation and tutorials)
- Installed Kuberikube on local machine.
- Set up a cluster on AWS EKS and went through tutorials to ensure its working.

Tarunianand Muruganandan:







- Began with working on the basic Django framework as a starting point to build an application.
- Have been working locally to understand the backend portion, at a bare bones level.
- Reading up on Spring and other RESTful web APIs that can be used for the project.
- I have been trying to make a presentable application, but still working on setting up the right environment. I have currently managed to run all migration and start the server, with some rendering issues left to be fixed.
- Learnt how submodules work on GitHub if we push a pre-existing repo into another repo.

```
Taxi_Application — python • python manage.py runserver — 80x24
^CTarunianands-MacBook-Pro:Taxi_Application taruni$ python manage.py migrate
Operations to perform:
  Apply all migrations: admin, auth, contenttypes, sessions, uber
Running migrations:
  Applying contenttypes.0001_initial... OK
  Applying auth.0001_initial... OK
  Applying admin.0001_initial... OK
  Applying admin.0002_logentry_remove_auto_add... OK
  Applying contenttypes.0002_remove_content_type_name... OK
  Applying auth.0002_alter_permission_name_max_length... OK
  Applying auth.0003_alter_user_email_max_length... OK
  Applying auth.0004_alter_user_username_opts... OK
  Applying auth.0005_alter_user_last_login_null... OK
  Applying auth.0006_require_contenttypes_0002... OK
  Applying auth.0007_alter_validators_add_error_messages... OK
  Applying auth.0008_alter_user_username_max_length... OK
  Applying auth.0009_alter_user_last_name_max_length... OK
  Applying sessions.0001_initial... OK
  Applying uber.0001_initial... OK
Tarunianands-MacBook-Pro:Taxi_Application taruni$ python manage.py runserver
Performing system checks...





System check identified no issues (0 silenced).
November 17, 2018 - 06:17:54
```

Github checkins





















Commits on Nov 16, 2018

Working pyspark to mongodb code AkritiKapur committed 39 minutes ago	 660f078	
db settings added AkritiKapur committed 3 hours ago	 4f502d6	
changed dynamo db to mongo db AkritiKapur committed 3 hours ago	 acf9e6b	







Commits on Nov 15, 2018

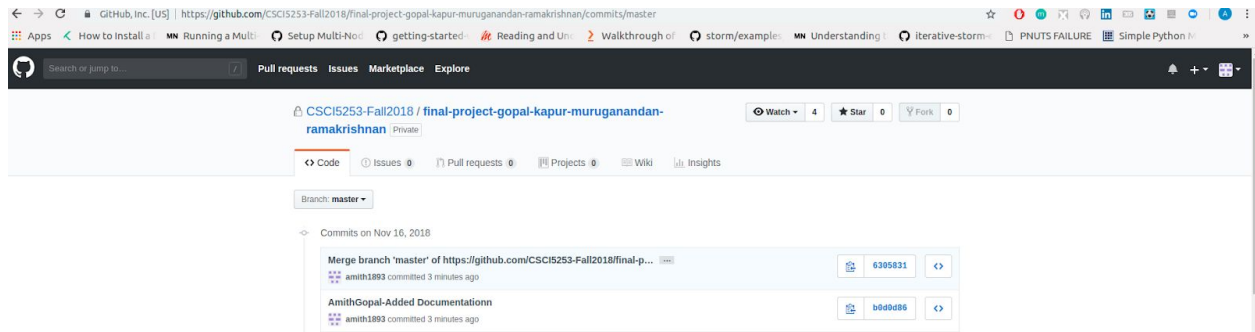
read csv to spark dataframe added AkritiKapur committed a day ago	 df0c0b8	
Init structure for pyspark dataframe and synamo db write AkritiKapur committed a day ago	 ea9c8d0	

Commits on Nov 14, 2018

Merge branch 'master' of github.com:CSCI5253-Fall2018/final-project-g... AkritiKapur committed 2 days ago	 443f095	
preprocessing-removed unused columns from dataset AkritiKapur committed 2 days ago	 0abeb4c	
updated requirements AkritiKapur committed 2 days ago	 4394ade	
added project settings file AkritiKapur committed 2 days ago	 80b55a6	
Delete ~/.lock.2014_Green_Taxi_Trip_Data.csv# AkritiKapur committed 2 days ago	Verified  f829c67	
environment, requirement file added AkritiKapur committed 2 days ago	 19d698e	
data folder added AkritiKapur committed 2 days ago	 ba78044	
preprocess lambda initial structure added AkritiKapur committed 2 days ago	 6e7d119	
Merge branch 'master' of github.com:CSCI5253-Fall2018/final-project-g... AkritiKapur committed 2 days ago	 f55bb5c	
initialize project structure AkritiKapur committed 2 days ago	 4824a71	

Commits on Nov 16, 2018

Added the application, not as a submodule Taruni-Anand committed 2 minutes ago	 1465632	
Added framework for application Taruni-Anand committed 28 minutes ago	 d186818	
Added framework for application Taruni-Anand committed 33 minutes ago	 37a83a8	



	Name	Tasks	Timeline
1	Akriti Kapur	<ul style="list-style-type: none"> • Kafka processing • Structure the queries • Creating test cases to produce “real-time” data • Simple queries to test Spark-dynamoDB 	<ul style="list-style-type: none"> • 19th Nov week • 26th Nov week • 26th Nov week
2	Amith Gopal	<ul style="list-style-type: none"> • Running custom application through Spark on EKS • Running different tests to test Spark on EKS further • Trying minikube as a fallback if we encounter issues in EKS • Integration of Application Layer with EKS/ECS 	<ul style="list-style-type: none"> • 19th Nov week • 19th Nov week • 26th Nov week • 26th Nov week
3	Sowmya	<ul style="list-style-type: none"> • Integration of Spark with Kubernetes • Integration of Application Layer with EKS/ECS • Documentation 	<ul style="list-style-type: none"> • 19th Nov week • 26th Nov week • 26th Nov week
4	Tarunianand	<ul style="list-style-type: none"> • Developing application layer, check how the backend would change to accomodate the rest of the project. • Trying different frameworks for the application (Spring, Flask etc.), finalize framework. • Integrating application layer with EKS/ECS • Creating test cases to produce “real-time” data • Documentation and Demonstration prep 	<ul style="list-style-type: none"> • 19th Nov week • 19th Nov week • 26th Nov week • 26th Nov week • 26th Nov week

Costs

Costs incurred

Estimated Costs for future use

Database Management

- We did not have to do a lot of cleaning on the database. There were few columns we did not need and hence removed it from the database.
- As our focus in the project is on scalability, we need to perform certain tasks on the kubernetes cluster. We may need additional columns depending on the tasks we need spark to perform on this database. These additional columns if needed, will be added to the database in the pre-processing step while using spark to write the data frame to MongoDB.

Challenges faced:

- Setting up EKS cluster was a challenge since there were a lot of steps to be followed and lot of prerequisites to be satisfied.
- Running spark on EKS. Had to figure out the permissions to send spark jobs from a local machine to a remote EKS cluster.
- Storing data from Spark into DynamoDB was difficult since we couldn't find any documentation online for saving pyspark dataframe to DynamoDB which is why we switched to MongoDB where the whole data storage pipeline is successfully constructed.
- We will be testing minikube as a fallback if there are any issues with EKS.

CHECKPOINT 2

Summary of the work done

- Got spark working on kubernetes clusters spun up on Minikube as well as AWS EKS
- Created an application using Django with Kafka integration
- Added map visualizations for taxi rides in New York
- Research on MongoDB integration with Kubernetes

Github Commits

Commits on Dec 4, 2018

Added maps using mapbox



Taruni-Anand committed 28 minutes ago



2cc5ee2



Commits on Dec 3, 2018

Readme for app updated, with steps to setup app and run kafka



AkritiKapur committed a day ago



151c78d



Readme for app updated, with steps to setup app and run kafka



AkritiKapur committed a day ago



b29a191



Readme for app updated, with steps to setup app and run kafka



AkritiKapur committed a day ago



cafc835



Kafka Consumer producer working code to add an entry to ride table



AkritiKapur committed a day ago



937bbf8



Commits on Dec 1, 2018

Kafka Consumer producer structure and test flow added



AkritiKapur committed 3 days ago



0e0f4df



Bulk create after every few objects added in the list to avoid crash



AkritiKapur committed 3 days ago



30d8b0a



Adding requirements for kafka - django-logpipe



AkritiKapur committed 4 days ago



ba57b80



Custom migration script for migrating data to django database



AkritiKapur committed 4 days ago



2a00772



Commits on Nov 27, 2018

create mongo model in new dbase



AkritiKapur committed 7 days ago



e5c25e2



django mongodb related requirements added



AkritiKapur committed 7 days ago



1d301a4



mongo db settings, model update



AkritiKapur committed 7 days ago



5297e4b



modified requirements for additional requirements



AkritiKapur committed 8 days ago



107cd6e



removing local database files, have to be replaced my mongodb settings



AkritiKapur committed 8 days ago



39fa9bb



	Name	Tasks	Timeline
1	Akriti Kapur	<ul style="list-style-type: none"> Connect Django/ Kafka with Kubernetes-Spark Work with Taruni on visualization for taxi rides nearby. 	Current week till December 12th
2	Amith Gopal	<ul style="list-style-type: none"> Currently reading up on how to connect to MongoDB from Kubernetes Cluster. Will be working with Sowmya on this one since it's an unexplored area and requires a lot of research. Will try running spark queries on MongoDB 	Current week till December 12th
3	Sowmya	<ul style="list-style-type: none"> Will be working with Amith to connect MongoDB from the Kubernetes Cluster. Will be working with Amith on this one since it's an unexplored area and requires a lot of research. Will be working on communication from Kubernetes cluster to the Django Application. 	Current week till December 12th
4	Tarunianand	<ul style="list-style-type: none"> Finalize the app configurations and make it user friendly, complete visualizations, make required migrations and connections Work with Akriti in connecting Django with Kubernetes-Spark 	Current week till December 12th

Akriti Kapur

- Integrating MongoDB with Django
- Creating Database models in Django - Django models correspond to database tables
- Migrating data to the model created through Django - A custom migration script to migrate rows from the old (cleaned) database to Django managed database. This was done so that the Django ORM can be used for visualizations later!
- Integrating Kafka with the application - Kafka would be used to put tasks (topics) on the queue like "find rides in the range {start time} and {end time} and in a longitude and latitude range"
- A producer and consumer test code added for Kafka
- Creating a [Readme](#) on instructions on how to get the application running with installation instructions.

Tarunianand Muruganandan

- Finalized Django as the framework to be used, weighing the pros and cons with other frameworks.
- Researched on some visualization tools and APIs to represent taxi travel on a map with Django.
- Finalized on using MapBox, got a basic map and navigation tool to work as a separate page on the application.
- Most of the progress was made locally, as they were run on local applications, that I built solely to test on. Once the rendering was satisfactory, it made its way to a commit.

Amith Gopal & Sowmya Ramakrishnan

Running Spark on EKS

The same steps as listed in the 'Integrating Spark and EKS' document were followed for successful creation of a Kubernetes cluster on AWS (EKS).

It was ensured that a user can list, edit, create and delete pods using the command

'kubectl auth can-i '*' '*' '

Installing Spark2.3 and running Spark jobs

- The Spark binary files were downloaded.
- The environment variables (JAVA_HOME - /usr/lib/jvm/java-8-openjdk-amd64/jre and SPARK_HOME - /home/kagura/Downloads/spark2.3) were set.
- The Kubernetes Cluster Master URL was obtained by running the command `kubectl cluster-info`.
- Kube-DNS (required to properly resolve cluster.local names) was ensured to be installed and functioning (running) properly using the command `'kubectl get pod -n kube-system'`
- To allow the driver pod to create pods and services (for the local machine to be given access to the Kubernetes cluster), two Cluster Role Bindings were created using the following commands:

```
kubectl create clusterrolebinding cluster-system-anonymous  
--clusterrole=cluster-admin --user=system:anonymous
```

```
kubectl create clusterrolebinding default-admin --clusterrole cluster-admin  
--serviceaccount=default:default
```

- Using the master URL and example jar (uploaded in a public S3 bucket), the spark-submit binary was used to execute the sample spark query (that calculates the value of Pi). The spark image was obtained from a Google cloud container registry.

```
spark-submit --master
k8s://https://962DF1838BF53D6B319A4ABA4913E410.yl4.us-east-1.eks.amazonaws.com
--deploy-mode cluster \
--name spark-pi \
--class org.apache.spark.examples.SparkPi \
--conf spark.executor.instances=5 \
--conf
spark.kubernetes.container.image=gcr.io/cloud-solutions-images/spark:v2.3.0-gcs
--conf spark.kubernetes.driver.pod.name=spark-pi-driver
http://sparkexamplejar.s3.amazonaws.com/spark-examples\_2.11-2.3.0.jar
```

- The job was monitored on the Kubernetes Dashboard and succeeded, with the logs having information and the calculated Pi values.

Running Spark on Minikube

- The following prerequisites were met: The Apache Spark 2.3 binary was downloaded, Docker, VirtualBox, Minikube and kubectl were installed on a Linux machine.
- Minikube was started locally with 8 gigs of memory and 4 cpu cores using the command 'minikube start --memory 8192 --cpus 4'
- A spark docker image was created locally using docker and the downloaded spark distribution. In the extracted spark folder, the following command was run: './bin/docker-image-tool.sh -m -t spark-docker build'. Running the command 'docker image ls' showed the available docker build.
- This image was then pushed to a public docker hub repository. To enable access to the repo, we first logged onto docker hub using the 'sudo docker login --username <username>' command. The following steps were involved in pushing the image to the repo:

```
sudo ./bin/docker-image-tool.sh -r sowmya2910/myrepo -t v2.3.0 build

sudo docker tag sowmya2910/myrepo/spark:v2.3.0 sowmya2910/myrepo-spark:v2.3.0

sudo docker push sowmya2910/myrepo-spark:v2.3.0
```

- To allow the driver pod to create pods and services (for the local machine to be given access to the Kubernetes cluster), two Cluster Role Bindings were created using the following commands:

```
kubectrl create clusterrolebinding cluster-system-anonymous  
--clusterrole=cluster-admin --user=system:anonymous  
  
kubectrl create clusterrolebinding default-admin --clusterrole cluster-admin  
--serviceaccount=default:default
```

- The spark job was submitted using the following command:

```
spark-submit \  
--master k8s://https://192.168.99.100:8843 \  
--deploy-mode cluster \  
--name spark-pi \  
--class org.apache.spark.examples.SparkPi \  
--conf spark.executor.instances=3 \  
--conf spark.kubernetes.container.image=sowmya2910/myrepo-spark:v2.3.0  
http://sparkexamplejar.s3.amazonaws.com/spark-examples\_2.11-2.3.0.jar
```

- The cluster was then deleted using the 'minikube stop' and 'minikube delete' commands in succession.
- The job was monitored on the Kubernetes Dashboard and succeeded, with the logs having information and the calculated Pi values.

Thus, we were able to run spark on kubernetes using both EKS as well as Minikube. The issues encountered, appropriate/consequent debugging performed and screenshots validating successful run are documented separately and uploaded in the project git repository.

Meetings

We have a channel through which we communicate. This channel is very active. We put down any notes about changes we are considering in our implementations or design here. Any contentions are discussed further. However, physically we have met the following times:

- We had a meeting to outline the tasks and divide our tasks into subtasks after checkpoint 1. All four of the teammates were a part of this meeting. This was a one-hour brainstorming meeting. We mainly discussed the division of tasks into subtasks and how to approach each subtask. For example, if someone was using Kafka what all libraries were he/she going to try, what was the approach, would this be part of Django or a different service etc.
- Amith and Sowmya had a meeting to get Spark working on Kubernetes. Both of us were constantly communicating and discussing about how to overcome the errors we were facing while deploying Spark on the Kubernetes cluster. We met on 4th December for about an hour to resolve some issues that were occurring in my(Amith) setup due to some discrepancies in the configuration. We also discussed about the future steps to be taken which involved how to go about connecting MongoDB from the Kubernetes cluster and how to communicate with the Django application.

Changes to the Proposal

There weren't any changes to the proposal. However, there are a few things (details) that weren't anticipated or mentioned as a part of the document. These are as follows:

- Visualizations - This was a stretch goal for us but we decided to include it in our project as this is the best way to showcase that the tasks like finding taxi rides etc are executed.
- Migrations - Django, the web framework works best with a Model-view-controller defined and since we are building a web application we decided this was the route to take and decided to have Django models.
- Using Minikube as an alternative to AWS EKS is being considered. Although we were able to run sample spark queries successfully on kubernetes clusters spun up on both the interfaces, we are leaning more toward Minikube primarily since integration of the MongoDB database with Kubernetes looks simpler with Minikube, and it is also a more rational choice, cost-wise.

Dataset Management

- We did not have any issues with scrubbing or cleaning data. However, it did take a considerable amount of time to perform migrations to Django. This was because a Django model needed to be created and a script needed to be written for copying the data from our original database to the new one. This was needed so that we could perform queries using the inbuilt Django ORM which would make things like visualizations, any data manipulations or get, post, put, delete calls easier.

Challenges

- Running spark on EKS. Had to figure out the permissions to send spark jobs from a local machine to a remote EKS cluster, and had to deal with the many issues that popped up, as outlined in the 'Work Done' section

- Running spark on Minikube. There were many intricacies involved and figuring out and debugging issues took quite some time and lots of research. Due to excellent community support (both Spark and Kubernetes) though, we were able to get it up and running.

Costs

Costs incurred

We used AWS EKS to spin up our Kubernetes cluster, and ran sample spark queries on it. The process of setting up Kubernetes using EKS is elaborate and required launch of a cluster VPC and worker nodes (CloudFormation) and the cluster plane (EKS). EKS costs at about \$0.20 per hour of use, and the AWS resources that it uses cost too. - t2.medium EC2 instances (\$0.05 * 3 per hour) and EBS volumes (\$0.10 per GB-month of provisioned storage). Setting up the cluster using EKS involves quite a bit of steps, and we could hence not terminate the cluster and its resources when not using them (that time was considerably less), this costed us about \$30 in the couple weeks that we used the resources. The starter account could not be used since it does not allow use of IAM, among everything else.

Estimated Costs for future use

We are trying to make things work with Minikube instead, but if we find that using EKS is simpler and opt for it instead, we might incur some cost. We can try keeping it at a minimum by terminating all resources and bringing them up again from time to time, which is again a tradeoff since it is way more tedious to do so.